

Abstract

This paper examines the predictive power of order book imbalance for near-term price movements using millisecond-resolution limit order book data from liquid U.S. equities. We analyze how imbalance, the relative difference between bid and ask side liquidity, predicts future returns at horizons ranging from 10 milliseconds to several seconds.

We find that order book imbalance contains economically and statistically significant predictive information at very short horizons, with Information Coefficients exceeding 0.25 at 10 milliseconds. This predictability decays rapidly and monotonically with forecast horizon, indicating that imbalance reflects transient information quickly incorporated into prices. The effect persists after controlling for bid-ask spread and short-term volatility, demonstrating that imbalance captures information distinct from standard microstructure variables.

Time-of-day analysis reveals regime-dependent patterns: imbalance is most informative during stable midday trading, exhibits weakest predictability at the open, and shows strongest price impact near the close. The regression coefficient β remains stable or increases across short horizons even as directional accuracy declines, suggesting imbalance reflects both informational content and mechanical price adjustment from liquidity consumption.

These findings establish order book imbalance as an ultra-short-term indicator of local supply-demand pressure rather than a persistent alpha signal. The results have direct implications for execution quality, adverse selection management, and short-horizon trading strategies in electronic markets.

Introduction

Modern electronic markets operate on millisecond-to-second timescales, where anticipating near-term price movements is essential for market makers and algorithmic traders managing inventory risk and execution quality. At these horizons, traditional predictors are largely irrelevant, and the microstructure of the limit order book becomes the dominant source of information.

Order book imbalance, the relative difference between liquidity on the bid and ask sides, is a natural candidate for prediction: asymmetric liquidity indicates the direction in which prices can move more easily. However, existing evidence is mixed, as imbalance is correlated with other microstructure variables, and its predictive power decays rapidly with horizon. Moreover, most studies focus on returns alone without linking statistical predictability to execution realities such as adverse selection.

This paper analyzes order book imbalance using millisecond-resolution data across multiple liquid U.S. equities. We examine how predictability decays from 10 milliseconds to several seconds, test robustness to standard microstructure controls, and analyze effects on execution outcomes. We find that imbalance contains significant predictive information at very short horizons, that this information decays systematically over time, and that it is not subsumed by spread or volatility. Additionally, it meaningfully affects adverse selection and post-trade price dynamics. These results clarify when and how order book imbalance functions as an informational signal rather than a purely mechanical feature of market microstructure.

Hypothesis

In electronic limit order markets, prices respond to the balance between supply and demand liquidity. The limit order book aggregates this balance by showing where liquidity is concentrated and how easily prices can move in either direction. Order book imbalance, therefore, summarizes short-term trading pressure and price impact asymmetry.

Imbalance can arise from three related mechanisms. Market makers manage inventory by shifting quotes asymmetrically, creating imbalances that precede price adjustments. Liquidity itself is often asymmetric, so trades on the thinner side of the book move prices more, causing prices to drift toward the side with less depth. Finally, large traders executing over time leave persistent footprints by repeatedly consuming liquidity on one side, generating sustained imbalance before prices fully adjust. These mechanisms imply that imbalance should predict near-term returns, but only over very short horizons, as the information is rapidly incorporated through execution and quote updates.

We therefore hypothesize that order book imbalance predicts short-horizon future returns, that this predictability decays monotonically as the horizon increases, that it is strongest during stable periods of continuous trading and weaker during the open and close, and that trades executed in the direction of imbalance suffer less adverse selection than trades executed against it.

Data

We use full-depth limit order book and trade data from the NASDAQ exchange for August 1, 2022, covering three liquid U.S. instruments: Apple Inc. (AAPL), the Invesco QQQ Trust (QQQ), and the SPDR S&P 500 ETF Trust (SPY). These securities exhibit continuous

trading, deep liquidity, and high message rates, making them suitable for short-horizon microstructure analysis.

The data consist of nanosecond-timestamped, sequenced message-level updates, including incremental order book events, periodic snapshots, top-of-book updates, and detailed execution records. This structure allows exact reconstruction of the order book through time and precise alignment between book state, trades, and subsequent price changes.

The order book and trade data are stored in fixed-layout binary structures packed without padding and validated via compile-time size checks.

The book is reconstructed from incremental messages using price–time priority keyed by order ID. The event stream includes order additions, deletions, replacements, quantity reductions, executions, hidden trades, and session events. Periodic snapshots provide consistency checks and recovery points.

For signal construction, we use explicit top-of-book structures containing the first three bid and ask levels (price and quantity). For execution analysis, trade records contain execution price and size as well as the resting order’s side, price, remaining quantity, and hidden status.

All messages carry nanosecond timestamps and monotonic sequence numbers. The book state is updated strictly in sequence order, and trades are aligned to the most recent book state with timestamp less than or equal to the execution time.

We restrict analysis to continuous trading hours (09:30–16:00 ET), excluding auctions, session transitions, crossed books, and observations with zero or undefined spreads.

The mid-price is defined as the average of the best bid and ask, and the quoted spread is their difference. Future returns over horizon Δ are computed as changes in the mid-price.

Order book imbalance is measured using both the top level (OBI_1) and the first three levels (OBI_3):

$$OBI_k = (\sum_{i=1}^k q_i^{bid}(t) - \sum_{i=1}^k q_i^{ask}(t)) / (\sum_{i=1}^k q_i^{bid}(t) + \sum_{i=1}^k q_i^{ask}(t)).$$

Volatility is proxied by a rolling standard deviation of mid-price returns, with absolute returns used as a robustness alternative. All variables are aligned to a common timestamp.

Methodology

We construct order book imbalance using the aggregated depth on the bid and ask sides of the limit order book. Let $q_i^{bid}(t)$ and $q_i^{ask}(t)$ denote the displayed quantity at the i -th price level on the bid and ask sides at time t , respectively. The k -level order book imbalance is defined as

$$OBI_k = (\sum_{i=1}^k q_i^{bid}(t) - \sum_{i=1}^k q_i^{ask}(t)) / (\sum_{i=1}^k q_i^{bid}(t) + \sum_{i=1}^k q_i^{ask}(t)).$$

We use $k=1$ (OBI_1) to capture immediate top-of-book pressure and $k=3$ (OBI_3) to capture slightly deeper liquidity conditions. Both measures are bounded in $[-1,1]$ with positive values indicating net bid-side dominance.

All analysis is conducted in event time rather than calendar time. Each observation corresponds to an order book update or trade event, rather than a fixed clock interval.

Event-time sampling is standard in high-frequency settings because it avoids the severe heteroskedasticity present in clock-time returns: periods of intense activity generate many observations with small price changes, while quiet periods generate few observations with larger

changes. Sampling in event time equalizes information arrival across observations and produces more stable error behavior.

For each event at time t , we evaluate future price changes over horizons $\Delta \in \{10, 100, 1000, 5000\}$ milliseconds, defined relative to the event timestamp.

Let $m(t)$ denote the mid-price at time t , defined as the average of the best bid and ask. For each event at time t , we define the future return over horizon Δ as

$$r(t+\Delta) = (m(t+\Delta) - m(t)) / m(t) \times 10,000,$$

expressed in basis points. The mid-price at $t+\Delta$ is obtained by locating the first book update with a timestamp greater than or equal to $t+\Delta$.

This definition captures the directional price movement faced by a trader who acts at time t and unwinds at horizon Δ .

We estimate the predictive relationship between imbalance and future returns using ordinary least squares:

$$r(t+\Delta) = \alpha + \beta \cdot OBI_k(t) + \varepsilon_t.$$

We report the slope coefficient β , its t-statistic, and the coefficient of determination R^2 for each horizon and each imbalance measure.

As a complementary nonparametric measure, we compute the information coefficient (IC), defined as the Spearman rank correlation between $OBI_k(t)$ and $r(t+\Delta)$.

To assess whether imbalance simply proxies for other microstructure variables, we estimate augmented regressions:

$$r(t+\Delta) = \alpha + \beta \cdot OBI_k(t) + \gamma \cdot \text{Spread}(t) + \delta \cdot \text{Vol}(t) + \varepsilon_t,$$

where Spread is the quoted bid–ask spread and Vol is a rolling proxy for short-term volatility. We examine coefficient stability, statistical significance, and changes in R^2 when controls are added.

We evaluate stability across market conditions by repeating all analyses across intraday regimes: the market open (09:30–10:00), midday (10:00–15:30), and close (15:30–16:00). This allows us to identify when imbalance is most informative and when its predictive power weakens.

To assess statistical stability under temporal dependence, we implement a moving-block bootstrap with one-second blocks. We resample contiguous blocks of observations with replacement and recompute coefficients across 300 bootstrap iterations. We report bootstrap means, standard deviations, and confidence intervals for β and IC.

We also compare results across OBI_1 and OBI_3 , across horizons, and across instruments to evaluate robustness.

Finally, we link statistical predictability to realized trading outcomes by conditioning adverse selection on whether trades are aligned or misaligned with the contemporaneous imbalance. We measure post-trade mid-price movements following executions to quantify how imbalance affects realized execution quality.

Results

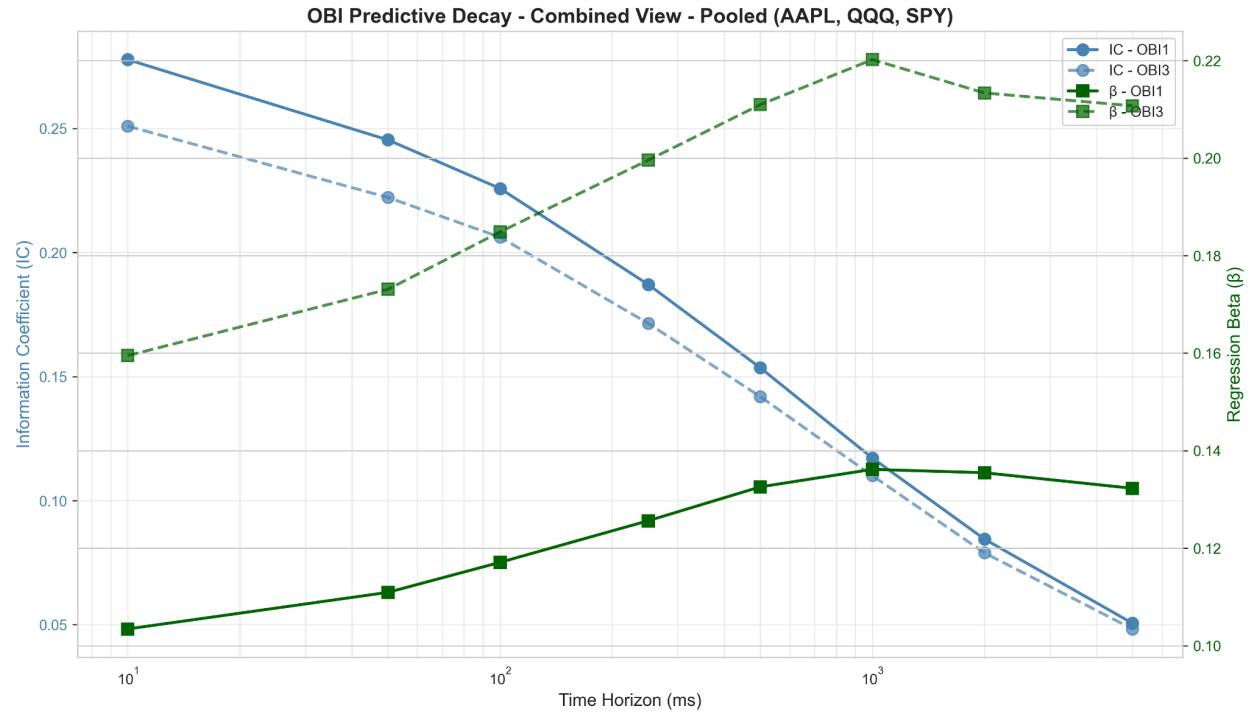


Figure 1: Predictive decay of order book imbalance.

Information coefficients (IC, left axis) and regression coefficients (β , right axis) are plotted against horizon Δ for OBI_1 and OBI_3 , pooled across AAPL, QQQ, and SPY. IC declines monotonically with horizon, while β rises at short horizons and then stabilizes.

Figure 2 shows how the predictive content of order book imbalance varies with the prediction horizon. The information coefficient (IC), which measures directional accuracy, declines monotonically as Δ increases, indicating that imbalance contains primarily short-lived information that is rapidly incorporated into prices.

In contrast, the regression coefficient β initially increases with horizon before flattening, suggesting that while imbalance becomes less directionally reliable, the conditional price response unfolds over slightly longer timescales. Together, these patterns indicate that imbalance reflects transient information combined with slower price adjustment driven by liquidity consumption and order flow dynamics.

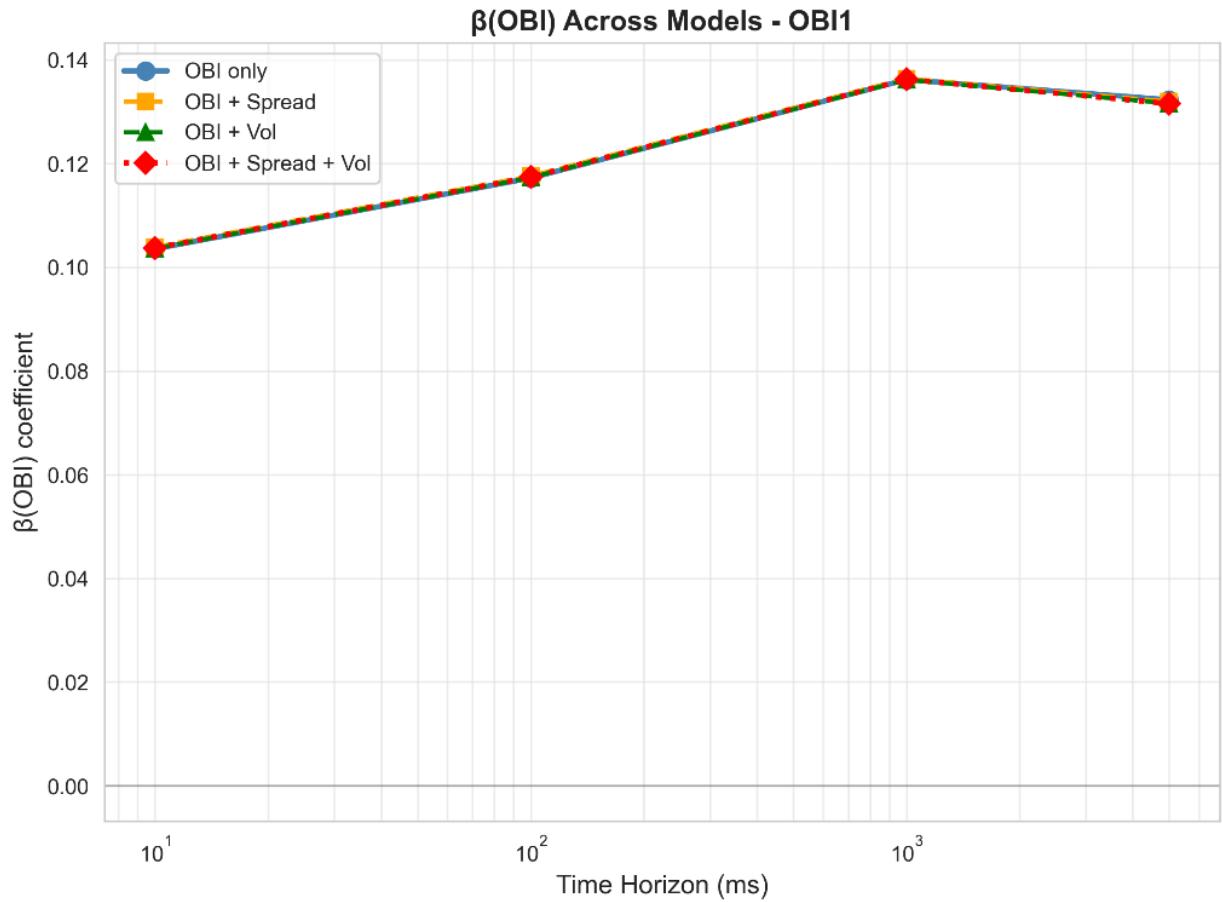


Figure 2: OBI coefficient across models and horizons (OBI_1)

Estimated β coefficients from regressions of future returns on OBI_1 across horizons Δ . Shown are univariate estimates and estimates controlling for bid–ask spread and short-term volatility. The OBI coefficient remains stable across specifications.

Table 1: OBI₁ Control Regressions

Δ	β only	β full	Retention	γ	δ	R^2
10	0.1035	0.1037	100.18%	-0.068	0.075	0.057
100	0.1172	0.1174	100.18%	-0.075	0.081	0.035
1000	0.1362	0.1362	100.02%	-0.047	0.080	0.009
5000	0.1323	0.1315	99.43%	0.031	0.165	0.002

OBI₁ regressions with liquidity and volatility controls. This table reports coefficients from regressions of future returns on OBI₁ alone and augmented with contemporaneous bid–ask spread and short-term volatility across horizons Δ . “ β retention” denotes the ratio of the OBI coefficient in the full specification to that in the univariate specification. The stability of β across models indicates that imbalance is not a proxy for spread or volatility.

Table 2: OBI₃ Control Regressions

Δ	β only	β full	Retention	γ	δ	R^2
10	0.1600	0.1611	100.71%	-0.070	0.072	0.051
100	0.1853	0.1866	100.68%	-0.077	0.077	0.033
1000	0.2209	0.2214	100.23%	-0.049	0.075	0.009
5000	0.2110	0.2083	98.75%	0.029	0.160	0.002

OBI₃ regressions with liquidity and volatility controls. This table reports coefficients from regressions of future returns on OBI₃ alone and augmented with contemporaneous bid–ask spread and short-term volatility across horizons Δ . “ β retention” denotes the ratio of the OBI coefficient in the full specification to that in the univariate specification. The persistence of β

across models indicates that OBI_3 captures information distinct from standard microstructure variables.

To assess whether order book imbalance merely proxies for other microstructure variables, we regress future returns on OBI along with contemporaneous bid–ask spread and short-term volatility. Figure 2 reports the estimated OBI_1 coefficients across horizons for the univariate specification and for models including controls.

Across all horizons, the OBI coefficient remains essentially unchanged after adding spread and volatility. For OBI_1 , coefficient retention exceeds 99% at all horizons, and for OBI_3 it exceeds 98%, indicating that the imbalance signal is not subsumed by either control. Both spread and volatility enter significantly with economically intuitive signs, and their inclusion modestly increases R^2 , but they do not attenuate the imbalance effect.

These results indicate that order book imbalance captures predictive information distinct from standard liquidity and volatility measures, rather than serving as a mechanical proxy for either. The predictive content of OBI therefore reflects informational asymmetries in the limit order book rather than simple transaction cost or risk effects.

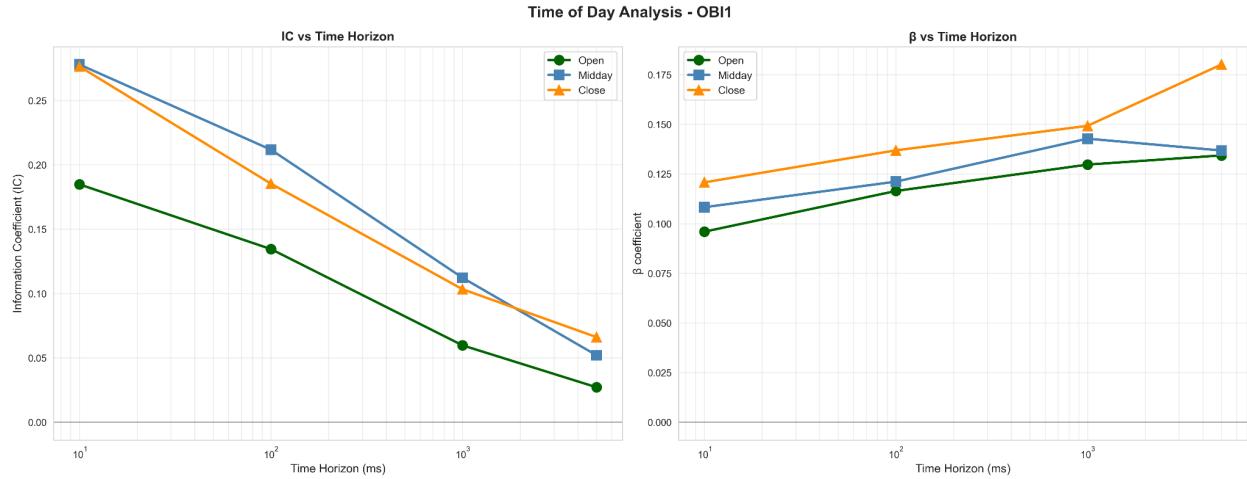


Figure 3: Time-of-day analysis for OBI₁.

The figure shows the Information Coefficient (left) and regression coefficient β (right) across horizons Δ for the open, midday, and close trading periods.

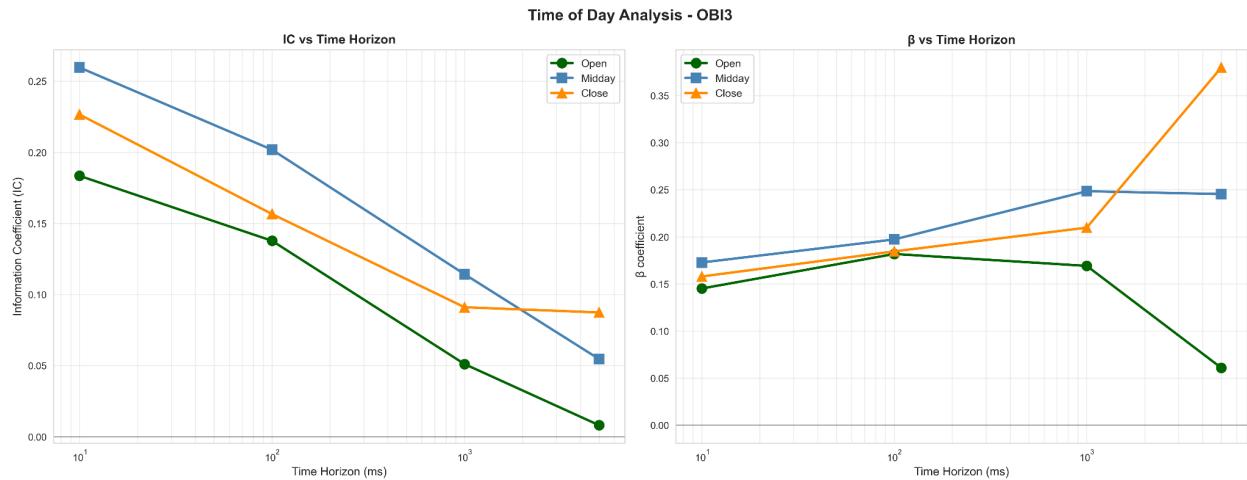


Figure 4: Time-of-day analysis for OBI₃.

The figure shows the Information Coefficient (left) and regression coefficient β (right) across horizons Δ for the open, midday, and close trading periods.

Table 3: OBI₁ Time-of-day Regression Results

Period	Δ	β	t-stat	IC	R ²	n
Open	10	0.0960	88.03	0.185	0.034	219205
Open	100	0.1164	63.55	0.134	0.018	219205
Open	1000	0.1297	28.01	0.060	0.004	219205
Open	5000	0.1344	12.73	0.027	0.001	219205
Midday	10	0.1083	308.02	0.278	0.077	1133472
Midday	100	0.1212	230.65	0.212	0.045	1133472
Midday	1000	0.1428	120.23	0.112	0.013	1133472
Midday	5000	0.1368	55.47	0.052	0.003	1133472
Close	10	0.1208	102.35	0.276	0.076	126628
Close	100	0.1369	67.10	0.185	0.034	126628
Close	1000	0.1492	36.96	0.103	0.011	126628
Close	5000	0.1801	23.58	0.066	0.004	126628

The table reports β , t-statistics, Information Coefficient (IC), R², and sample sizes across horizons Δ for the open, midday, and close trading periods.

Table 4: OBI₃ Time-of-day Regression Results

Period	Δ	β	t-stat	IC	R ²	n

Open	10	0.1452	87.39	0.183	0.034	219205
Open	100	0.1818	65.14	0.138	0.019	219205
Open	1000	0.1692	23.96	0.051	0.003	219205
Open	5000	0.0609	3.79	0.008	0.000	219205
Midday	10	0.1728	286.28	0.260	0.067	1133472
Midday	100	0.1973	219.43	0.202	0.041	1133472
Midday	1000	0.2485	122.52	0.114	0.013	1133472
Midday	5000	0.2455	58.27	0.055	0.003	1133472
Close	10	0.1580	82.79	0.227	0.051	126628
Close	100	0.1846	56.43	0.157	0.025	126628
Close	1000	0.2098	32.53	0.091	0.008	126628
Close	5000	0.3799	31.22	0.087	0.008	126628

The table reports β , t-statistics, Information Coefficient (IC), R^2 , and sample sizes across horizons Δ for the open, midday, and close trading periods.

Figure 3 and Figure 4 report the predictive performance of OBI across trading regimes defined by time of day. We split the sample into the open, midday, and close and estimate predictive relationships separately within each regime.

Across both OBI_1 and OBI_3 , predictability as measured by the Information Coefficient declines rapidly with the forecast horizon in all regimes. However, the level and persistence of predictability vary systematically across the trading day. The midday period exhibits the highest short-horizon IC, while the close retains relatively more predictive power at longer horizons. The open consistently displays the weakest predictability, particularly beyond 1 second.

The behavior of the regression coefficient β differs from that of IC. While IC decays monotonically with horizon, β is stable or increasing across horizons, especially during the close. For OBI₃, the close exhibits a pronounced increase in β at longer horizons, indicating elevated sensitivity of prices to imbalance late in the trading day.

Tables 3 and 4 confirm that these patterns are statistically significant. Coefficients are precisely estimated across all regimes and horizons, with large t-statistics and economically meaningful magnitudes. The results indicate that order book imbalance is most informative during periods of high liquidity and active trading, and that market conditions near the close amplify the price impact of imbalance even as directional predictability weakens.

Overall, the time-of-day analysis shows that imbalance predictability is not uniform across regimes: it is strongest intraday at very short horizons, decays quickly in all periods, and is associated with larger price responses near the close.

Table 5: Summary of Main Results

Signal	Δ (ms)	Regime	β	t-stat	IC	R²	Retention	Sig
OBI ₁	10	Pooled	0.104	284	0.28	0.057	1.002	✓
OBI ₁	100	Pooled	0.117	222	0.23	0.035	1.002	✓
OBI ₁	1000	Pooled	0.136	118	0.12	0.009	1.000	✓
OBI ₁	5000	Pooled	0.132	54	0.05	0.002	0.994	✓
OBI ₃	10	Pooled	0.161	263	0.25	0.051	1.007	✓
OBI ₃	100	Pooled	0.187	211	0.21	0.033	1.007	✓
OBI ₃	1000	Pooled	0.221	115	0.11	0.009	1.002	✓

OBI ₃	5000	Pooled	0.208	52	0.05	0.002	0.988	✓
OBI ₁	100	Open	0.116	64	0.13	0.018	—	✓
OBI ₁	100	Midday	0.121	231	0.21	0.045	—	✓
OBI ₁	100	Close	0.137	67	0.19	0.034	—	✓
OBI ₃	100	Open	0.182	65	0.14	0.019	—	✓
OBI ₃	100	Midday	0.197	219	0.20	0.041	—	✓
OBI ₃	100	Close	0.185	56	0.16	0.025	—	✓

Summary of predictive performance and robustness across signals, horizons, and regimes. The table reports effect sizes (β), statistical significance (t-stat), Information Coefficient (IC), R^2 , and robustness to microstructure controls (β retention). Significance is indicated at the 5% level.

Table 5 summarizes the main empirical findings across signals, horizons, and regimes.

Both OBI₁ and OBI₃ exhibit strong short-horizon predictability that decays rapidly with the forecast horizon. Predictability is highest during midday trading and weakest at the open, while price impact as measured by β is largest near the close. Across all specifications, imbalance remains statistically and economically significant after controlling for spread and volatility.

Execution Implications and Limitations

Although this study does not directly measure execution outcomes such as slippage, fill probability, or queue position, the results have clear implications for execution and market making.

The strong short-horizon predictability of order book imbalance implies that trading *against* imbalance exposes liquidity takers to adverse short-term price movements, while

conditioning on imbalance can reduce adverse selection risk. However, the rapid decay of predictability shows that imbalance is not a persistent alpha signal but an ultra-short-term indicator of local price pressure.

The time-of-day results further indicate that imbalance is most informative during stable continuous trading periods (midday and close) and weakest near the open, when price discovery dominates and the book is noisier.

Overall, order book imbalance should be viewed primarily as an execution and risk-management signal rather than as a medium-horizon return predictor.

The analysis is based on a single trading day, which limits inference about robustness across market regimes and prevents true out-of-sample testing. Execution outcomes such as realized slippage, fill rates, and queue position are not directly observed, so execution implications are inferred rather than measured.

In addition, while the regressions control for spread and volatility, other microstructure variables may interact with imbalance and are not explicitly modeled. Finally, the focus on highly liquid U.S. equities limits generalizability to other assets and market conditions.

Conclusion

This paper studies order book imbalance as a short-horizon predictive signal using millisecond-resolution data for liquid U.S. equities. The analysis shows that imbalance contains

economically and statistically significant information about near-term price movements, but this information decays rapidly with the forecast horizon.

Predictive accuracy declines monotonically from 10 milliseconds to several seconds, while the price response coefficient remains stable or increases, suggesting that imbalance captures both transient information and mechanical price impact from liquidity consumption. The imbalance effect is not subsumed by spread or volatility controls, and time-of-day analysis reveals it is most informative during midday trading, weakest at the open, and shows strongest price impact near the close.

These results imply that order book imbalance is not a persistent alpha signal but an ultra-short-term indicator of local supply–demand pressure, most useful for execution, inventory management, and short-horizon risk control. Conditioning trades on imbalance can reduce adverse selection, but its usefulness decays quickly and depends strongly on market conditions.

The analysis has several limitations: it covers a single trading day and small set of instruments, limiting generalizability; execution outcomes like slippage and fill rates are inferred rather than measured; and extending across longer periods and different market environments would strengthen inference.

Overall, the findings clarify that order book imbalance functions as an informational signal primarily at very short horizons, is robust to standard microstructure controls, varies across trading regimes, and is best interpreted as a tool for managing short-term price pressure rather than as a medium-horizon return predictor.