

Introduction

Modern electronic markets operate on millisecond-to-second timescales, where anticipating near-term price movements is essential for market makers and algorithmic traders managing inventory risk and execution quality. At these horizons, traditional predictors are largely irrelevant, and the microstructure of the limit order book becomes the dominant source of information.

Order book imbalance, the relative difference between liquidity on the bid and ask sides, is a natural candidate for prediction: asymmetric liquidity indicates the direction in which prices can move more easily. However, existing evidence is mixed, as imbalance is correlated with other microstructure variables, and its predictive power decays rapidly with horizon. Moreover, most studies focus on returns alone without linking statistical predictability to execution realities such as adverse selection.

This paper analyzes order book imbalance using millisecond-resolution data across multiple liquid U.S. equities. We examine how predictability decays from 10 milliseconds to several seconds, test robustness to standard microstructure controls, and analyze effects on execution outcomes. We find that imbalance contains significant predictive information at very short horizons, that this information decays systematically over time, and that it is not subsumed by spread or volatility. Additionally, it meaningfully affects adverse selection and post-trade price dynamics. These results clarify when and how order book imbalance functions as an informational signal rather than a purely mechanical feature of market microstructure.

Hypothesis

In electronic limit order markets, prices respond to the balance between supply and demand liquidity. The limit order book aggregates this balance by showing where liquidity is concentrated and how easily prices can move in either direction. Order book imbalance, therefore, summarizes short-term trading pressure and price impact asymmetry.

Imbalance can arise from three related mechanisms. Market makers manage inventory by shifting quotes asymmetrically, creating imbalances that precede price adjustments. Liquidity itself is often asymmetric, so trades on the thinner side of the book move prices more, causing prices to drift toward the side with less depth. Finally, large traders executing over time leave persistent footprints by repeatedly consuming liquidity on one side, generating sustained imbalance before prices fully adjust. These mechanisms imply that imbalance should predict near-term returns, but only over very short horizons, as the information is rapidly incorporated through execution and quote updates.

We therefore hypothesize that order book imbalance predicts short-horizon future returns, that this predictability decays monotonically as the horizon increases, that it is strongest during stable periods of continuous trading and weaker during the open and close, and that trades executed in the direction of imbalance suffer less adverse selection than trades executed against it.

Data

We use full-depth limit order book and trade data from the NASDAQ exchange for August 1, 2022, covering three liquid U.S. instruments: Apple Inc. (AAPL), the Invesco QQQ Trust (QQQ), and the SPDR S&P 500 ETF Trust (SPY). These securities exhibit continuous

trading, deep liquidity, and high message rates, making them suitable for short-horizon microstructure analysis.

The data consist of nanosecond-timestamped, sequenced message-level updates, including incremental order book events, periodic snapshots, top-of-book updates, and detailed execution records. This structure allows exact reconstruction of the order book through time and precise alignment between book state, trades, and subsequent price changes.

The order book and trade data are stored in fixed-layout binary structures packed without padding and validated via compile-time size checks.

The book is reconstructed from incremental messages using price–time priority keyed by order ID. The event stream includes order additions, deletions, replacements, quantity reductions, executions, hidden trades, and session events. Periodic snapshots provide consistency checks and recovery points.

For signal construction, we use explicit top-of-book structures containing the first three bid and ask levels (price and quantity). For execution analysis, trade records contain execution price and size as well as the resting order’s side, price, remaining quantity, and hidden status.

All messages carry nanosecond timestamps and monotonic sequence numbers. The book state is updated strictly in sequence order, and trades are aligned to the most recent book state with timestamp less than or equal to the execution time.

We restrict analysis to continuous trading hours (09:30–16:00 ET), excluding auctions, session transitions, crossed books, and observations with zero or undefined spreads.

The mid-price is defined as the average of the best bid and ask, and the quoted spread is their difference. Future returns over horizon Δ are computed as changes in the mid-price.

Order book imbalance is measured using both the top level (OBI_1) and the first three levels (OBI_3):

$$OBI_k = (\sum_{i=1}^k q_i^{bid}(t) - \sum_{i=1}^k q_i^{ask}(t)) / (\sum_{i=1}^k q_i^{bid}(t) + \sum_{i=1}^k q_i^{ask}(t)).$$

Volatility is proxied by a rolling standard deviation of mid-price returns, with absolute returns used as a robustness alternative. All variables are aligned to a common timestamp.

Methodology

We construct order book imbalance using the aggregated depth on the bid and ask sides of the limit order book. Let $q_i^{bid}(t)$ and $q_i^{ask}(t)$ denote the displayed quantity at the i -th price level on the bid and ask sides at time t , respectively. The k -level order book imbalance is defined as

$$OBI_k = (\sum_{i=1}^k q_i^{bid}(t) - \sum_{i=1}^k q_i^{ask}(t)) / (\sum_{i=1}^k q_i^{bid}(t) + \sum_{i=1}^k q_i^{ask}(t)).$$

We use $k=1$ (OBI_1) to capture immediate top-of-book pressure and $k=3$ (OBI_3) to capture slightly deeper liquidity conditions. Both measures are bounded in $[-1,1]$ with positive values indicating net bid-side dominance.

All analysis is conducted in event time rather than calendar time. Each observation corresponds to an order book update or trade event, rather than a fixed clock interval.

Event-time sampling is standard in high-frequency settings because it avoids the severe heteroskedasticity present in clock-time returns: periods of intense activity generate many observations with small price changes, while quiet periods generate few observations with larger

changes. Sampling in event time equalizes information arrival across observations and produces more stable error behavior.

For each event at time t , we evaluate future price changes over horizons $\Delta \in \{10, 100, 1000, 5000\}$ milliseconds, defined relative to the event timestamp.

Let $m(t)$ denote the mid-price at time t , defined as the average of the best bid and ask. For each event at time t , we define the future return over horizon Δ as

$$r(t+\Delta) = (m(t+\Delta) - m(t)) / m(t) \times 10,000,$$

expressed in basis points. The mid-price at $t+\Delta$ is obtained by locating the first book update with a timestamp greater than or equal to $t+\Delta$.

This definition captures the directional price movement faced by a trader who acts at time t and unwinds at horizon Δ .

We estimate the predictive relationship between imbalance and future returns using ordinary least squares:

$$r(t+\Delta) = \alpha + \beta \cdot OBI_k(t) + \varepsilon_t.$$

We report the slope coefficient β , its t-statistic, and the coefficient of determination R^2 for each horizon and each imbalance measure.

As a complementary nonparametric measure, we compute the information coefficient (IC), defined as the Spearman rank correlation between $OBI_k(t)$ and $r(t+\Delta)$.

To assess whether imbalance simply proxies for other microstructure variables, we estimate augmented regressions:

$$r(t+\Delta) = \alpha + \beta \cdot OBI_k(t) + \gamma \cdot \text{Spread}(t) + \delta \cdot \text{Vol}(t) + \varepsilon_t,$$

where Spread is the quoted bid–ask spread and Vol is a rolling proxy for short-term volatility. We examine coefficient stability, statistical significance, and changes in R^2 when controls are added.

We evaluate stability across market conditions by repeating all analyses across intraday regimes: the market open (09:30–10:00), midday (10:00–15:30), and close (15:30–16:00). This allows us to identify when imbalance is most informative and when its predictive power weakens.

To assess statistical stability under temporal dependence, we implement a moving-block bootstrap with one-second blocks. We resample contiguous blocks of observations with replacement and recompute coefficients across 300 bootstrap iterations. We report bootstrap means, standard deviations, and confidence intervals for β and IC.

We also compare results across OBI_1 and OBI_3 , across horizons, and across instruments to evaluate robustness.

Finally, we link statistical predictability to realized trading outcomes by conditioning adverse selection on whether trades are aligned or misaligned with the contemporaneous imbalance. We measure post-trade mid-price movements following executions to quantify how imbalance affects realized execution quality.